

# Task C.2 – Data Processing 1

Anh Vu Le

104653505

File: data\_processing.py

## Contents

(a) Start & End Dates .....	2
(b) NaN Handling .....	3
(c) Train/Test Split Methods .....	4
(d) Local Storage & Caching .....	5
(e) Feature Scaling & Scaler Storage .....	7
Conclusion.....	8

## (a) Start & End Dates

**Code (lines ~87–95):**

# Requirement (a): Write a function to load and process a dataset with multiple features  
a. This function will allow you to specify the start date and the end date for the whole dataset as inputs

```
78 #-----  
79 # REQUIREMENT (a): Allow user to specify start date and end date for whole dataset  
80 # Input validation for date parameters  
81 #-----  
82  
83 try:  
84     # Validate date format by attempting to parse them  
85     pd.to_datetime(start_date)  
86     pd.to_datetime(end_date)  
87     print(f"✓ Requirement (a): Date range specified - {start_date} to {end_date}")  
88 except:  
89     raise ValueError("start_date and end_date must be in 'YYYY-MM-DD' format")  
90  
91 if pd.to_datetime(start_date) >= pd.to_datetime(end_date):  
92     raise ValueError("start_date must be earlier than end_date")  
93
```

### Explanation

- `pd.to_datetime(start_date/end_date)` parses the two strings into real dates and will throw if the format is wrong (guards against bad inputs).
- Immediately validating dates prevents downstream silent failures in splitting/scaling that assume a valid window.

## (b) NaN Handling

Code (lines ~203–241; key ops at 218, 223, 228):

# Requirement (b): This function will allow you to deal with the NaN issue in the data.

```
198 #-----
199 # REQUIREMENT (b): Deal with NaN issue in the data
200 # Handle NaN (missing) values properly
201 #-----
202
203 print(f"✓ Requirement (b): Handling NaN values in the data")
204
205 # Check for NaN values in our selected features
206 nan_counts = feature_data.isnull().sum()
207 total_nans = nan_counts.sum()
208
209 if total_nans > 0:
210     print(f"Found {total_nans} NaN values:")
211     for col, count in nan_counts.items():
212         if count > 0:
213             print(f"  {col}: {count} NaN values")
214
215     # Strategy 1: Forward fill (use previous day's value)
216     # This is reasonable for stock prices as they tend to be continuous
217     # method='ffill' means forward fill - propagate last valid observation forward
218     feature_data_clean = feature_data.fillna(method='ffill')
219
220     # Strategy 2: Backward fill for any remaining NaN at the beginning
221     # This handles cases where the first few rows have NaN values
222     # method='bfill' means backward fill - use next valid observation to fill gap
223     feature_data_clean = feature_data_clean.fillna(method='bfill')
224
225     # Strategy 3: Drop any remaining rows with NaN (as last resort)
226     # If there are still NaN values after forward and backward fill, remove those rows
227     rows_before_drop = len(feature_data_clean)
228     feature_data_clean = feature_data_clean.dropna()
229     rows_after_drop = len(feature_data_clean)
230
```

### Explanation

- Logs per-column NaN counts, then applies **forward fill** → **backward fill** → **final drop** (only if something is still missing).
- This ordering is a common, conservative pattern for financial time series: carry the last valid price forward, patch early gaps, and finally drop any stubborn rows.
- The prints (“dropped X rows”, “After cleaning: 0 NaN”) are good, concrete evidence for your report.

## (c) Train/Test Split Methods

**Code (lines ~247–316):**

# Requirement (c): This function will also allow you to use different methods to split the data into train/test data; e.g. you can split it according to some specified ratio of train/test and you can specify to split it by date or randomly.

```
242 #-----
243 # REQUIREMENT (c): Use different methods to split data into train/test
244 # Split data into training and testing sets using flexible methods
245 #-----
246
247 print(f"✓ Requirement (c): Splitting data using '{split_method}' method")
248
249 if split_method == 'ratio':
250     # Method 1: Split by ratio - first X% for training, remaining for testing
251     # This maintains temporal order which is important for time series data
252     # We don't shuffle the data because time order matters in stock prediction
253     split_index = int(len(feature_data_clean) * split_value)
254
255     raw_train_data = feature_data_clean.iloc[:split_index].copy()
256     raw_test_data = feature_data_clean.iloc[split_index:].copy()
257
258     print(f"Ratio split ({split_value:.1%}):")
259     print(f"  Training: {len(raw_train_data)} samples ({raw_train_data.index[0]} to {raw_train_data.index[-1]})")
260     print(f"  Testing: {len(raw_test_data)} samples ({raw_test_data.index[0]} to {raw_test_data.index[-1]})")
261
262 elif split_method == 'date':
263     # Method 2: Split by specific date - all data before split_value for training
264     # This is useful when you want to test on a specific time period
265     # For example, train on 2020-2023 data, test on 2023-2024 data
266     split_date = pd.to_datetime(split_value)
267
268     raw_train_data = feature_data_clean[feature_data_clean.index < split_date].copy()
269     raw_test_data = feature_data_clean[feature_data_clean.index >= split_date].copy()
270
```

### Explanation

- **Ratio:** splits by proportion using `.iloc[:k]` (keeps temporal order intact).
- **Date:** converts the cutoff using `pd.to_datetime(split_value)` and filters by index `<` / `>=` that date (clear temporal separation).
- **Random:** samples indices with `train_test_split(..., random_state=42)` (reproducible), then **sorts** inside each subset so sequences remain chronological for sequence models.
- Immediately after this block your code prints sizes and date ranges of train/test, which is perfect to paste into the appendix as verification.

## (d) Local Storage & Caching

**Code (lines ~99–154):**

# Requirement (d): This function will have the option to allow you to store the downloaded data on your local machine for future uses and to load the data locally to save time

```
94 #-----
95 # REQUIREMENT (d): Store downloaded data locally and load from cache
96 # Setup caching system to avoid repeated downloads
97 #-----
98
99 print(f"✓ Requirement (d): Setting up local data caching")
100
101 # Create cache directory if it doesn't exist
102 # This allows us to store downloaded data locally for future use
103 if not os.path.exists(cache_dir):
104     os.makedirs(cache_dir)
105     print(f"Created cache directory: {cache_dir}")
106
107 # Generate unique cache key based on ticker and date range
108 # This ensures we can cache different datasets separately
109 cache_key = f"{ticker}_{start_date}_{end_date}"
110 cache_csv_path = os.path.join(cache_dir, f"{cache_key}.csv")
111 cache_meta_path = os.path.join(cache_dir, f"{cache_key}_meta.json")
112
113 # Check if cached data exists and load it
114 if os.path.exists(cache_csv_path) and os.path.exists(cache_meta_path):
115     print(f"✓ Loading from cache: {cache_csv_path}")
116     # Load the CSV file with proper date parsing
117     # index_col=0 means first column (Date) becomes the index
118     # parse_dates=True converts the index to datetime objects
119     data = pd.read_csv(cache_csv_path, index_col=0, parse_dates=True)
120
121     # Load metadata to verify cache validity
122     with open(cache_meta_path, 'r') as f:
123         cache_metadata = json.load(f)
124         print(f"Cache created: {cache_metadata['cache_date']}")
125
126 else:
127     print(f"✓ Downloading fresh data for {ticker} from {start_date} to {end_date}")
128
129     # Download data using yfinance
130     # yfinance is more reliable than pandas_datareader for Yahoo Finance data
131     data = yf.download(ticker, start=start_date, end=end_date)
132
133     # Handle potential multi-level column structure from yfinance
134     # Sometimes yfinance returns MultiIndex columns, we want simple column names
135     if isinstance(data.columns, pd.MultiIndex):
136         # Get the first level of column names (the actual feature names)
137         data.columns = data.columns.get_level_values(0)
138
139     # Save to cache for future use
140     data.to_csv(cache_csv_path)
141     print(f"✓ Data cached to: {cache_csv_path}")
```

**Explanation**

- Checks for a cached CSV + metadata JSON; if present, loads them with `parse_dates=True` so the index is datetime.
- If missing, downloads via `yfinance` (line ~130), flattens MultiIndex columns when needed, then writes both the CSV and a metadata JSON (ticker, range, columns, shape, cache time).
- This meets the “store & load locally” requirement, speeds up experiments, and gives provenance via JSON.

## (e) Feature Scaling & Scaler Storage

**Code (lines ~324–395+):**

# Requirement (e): This function will also allow you to have an option to scale your feature columns and store the scalers in a data structure to allow future access to these scalers

```
315 #-----
316 # REQUIREMENT (e): Scale feature columns and store scalers
317 # Apply feature scaling with proper handling to prevent data leakage
318 # This addresses ISSUE #2 mentioned in v0.1 comments
319 #-----
320
321 scalers = {}
322
323 if scale_features:
324     print(f"✓ Requirement (e): Applying MinMax scaling and storing scalers")
325     print(f"Scaling mode: {scale_mode}")
326     print("IMPORTANT: Fitting scalers on training data only to prevent data leakage")
327
328     # Initialize scaled datasets as copies of raw data
329     train_data = raw_train_data.copy()
330     test_data = raw_test_data.copy()
331
332     if scale_mode == 'all_features':
333         # Task C.5: Scale all features together using one scaler
334         scaler = MinMaxScaler(feature_range=(0, 1))
335
336         # Fit on the entire training dataframe
337         scaler.fit(raw_train_data[normalized_features])
338
339         # Transform both train and test data
340         train_data[normalized_features] = scaler.transform(raw_train_data[normalized_features])
341         test_data[normalized_features] = scaler.transform(raw_test_data[normalized_features])
342
343         # Store the single scaler
344         scalers['all'] = scaler
345         print("✓ Applied a single scaler to all features.")
346
347     if scale_mode == 'per_feature':
348         # Task C.6: Scale each feature column individually using one scaler per column
349         # This addresses ISSUE #2 mentioned in v0.1 comments
350         # Initialize a dict to store scalers for each feature column
351         feature_scalers = {}
352
353         # Iterate over each feature column and fit a scaler on the training data
354         for feature in normalized_features:
355             # Reshape the training data for each feature to (n, 1)
356             feature_train_data = raw_train_data[feature].values.reshape(-1, 1)
357             # Reshape the test data for each feature to (n, 1)
358             feature_test_data = raw_test_data[feature].values.reshape(-1, 1)
359
360             # Fit a MinMaxScaler on the training data for this feature
361             scaler = MinMaxScaler(feature_range=(0, 1))
362             scaler.fit(feature_train_data)
363
364             # Transform both train and test data for this feature
365             feature_train_data = scaler.transform(feature_train_data)
366             feature_test_data = scaler.transform(feature_test_data)
367
368             # Store the scaler for this feature
369             feature_scalers[feature] = scaler
370
371         # Store the dict of scalers
372         scalers['per_feature'] = feature_scalers
373
374         # Transform the training and test data using the per-feature scalers
375         for feature in normalized_features:
376             train_data[feature] = feature_train_data
377             test_data[feature] = feature_test_data
378
379         # Print a message indicating that per-feature scaling was applied
380         print("✓ Applied per-feature scaling to all features.")
381
382     # Print a message indicating that feature scaling was applied
383     print("✓ Applied feature scaling to all features.")
384
385     # Persist the scalers to a file
386     # This addresses ISSUE #2 mentioned in v0.1 comments
387     # Initialize a file path to store the scalers
388     scalers_file_path = os.path.join(directory, '_scalers.pkl')
389
390     # Persist the scalers to the file
391     pickle.dump(scalers, open(scalers_file_path, 'wb'))
392
393     # Print a message indicating that scalers were persisted
394     print(f"✓ Persisted scalers to {scalers_file_path}")
395
396     # Return the scaled training and test data
397     return train_data, test_data
398
399 # End of function
400
401 # Requirement (e) is complete
402
403 # Print a message indicating that requirement (e) is complete
404 print("✓ Requirement (e) is complete")
```

### Explanation

- Supports **two modes**:
  - **all\_features**: one MinMaxScaler across the full feature matrix.
  - **per\_feature**: one scaler per column (each column reshaped to 2D with `.values.reshape(-1,1)` because scikit-learn scalers expect `(n, 1)`).
- Every scaler is **fit on training data only** (no leakage), then applied to both train and test.
- Fitted scalers are stored in the scalers dict and also persisted to `..._scalers.pkl` (a few lines below this block in your file) so later stages can reuse them (e.g., inverse transform during inference).

## Conclusion

The enhanced data processing function:

Fixes all major issues in v0.1,

Implements requirements (a)–(e),

Provides a reproducible and extensible pipeline for later tasks.

This lays the groundwork for improved model training in Task C.3 and beyond.