

# COS30018 - Option C - Task 7 Report: Sentiment-Based Prediction

Anh Vu Le  
104653505

## I. Introduction and Executive Summary

### 1.1 Project Goal and Transition

Task C.7 marks a significant shift in the FinTech101 project from **Regression** (predicting the exact future price) to **Classification** (predicting the next-day price movement: UP or DOWN). The core goal is to build a robust model for the stock ticker **CBA.AX** (Commonwealth Bank of Australia) by integrating two distinct data modalities: **Technical Indicators** (from historical price data) and **Sentiment Features** (from external news sources).

### 1.2 Key Enhancement and Flow

The key enhancement of this task is the incorporation of market psychology through news sentiment. The end-to-end pipeline is orchestrated by `task7_runner.py` and follows this main flow:

1. **Data Collection:** Scrape news articles using `web_scraper.py`.
2. **Sentiment Analysis:** Generate daily sentiment scores (mean, volatility).
3. **Feature Engineering:** Combine 14 Technical Indicators + 2 Sentiment Features using `feature_builder.py`.
4. **Modelling:** Train 21 different classification models (XGBoost, Gradient Boosting, etc.).
5. **Evaluation:** Compare performance against a technical-only baseline to quantify sentiment's value.

### 1.3 Main Finding

A comprehensive comparative study showed that sentiment features add significant predictive power. The best-performing model, **Gradient Boosting (Sentiment-Only Features)**, achieved an **F1 Score of 68.8%**. This result demonstrated a substantial **+4.8% improvement** in F1 Score compared to the best Technical-Only Baseline (F1 = 65.7%), confirming the hypothesis that news sentiment captures valuable forward-looking information.

## Section II. Data Collection & Preprocessing

This section details the complete methodology for acquiring, cleaning, and synchronizing the two distinct data streams required for the analysis: (1) historical stock prices and (2) textual news data.

### Chunk 1: Stock Data Collection (The Temporal Backbone)

**Flow:** The entire project's timeline is built upon the stock market's trading calendar. This data stream acts as the "temporal backbone" for the project. The flow is initiated in the `task7_runner.py` script. This script calls the `yfinance` library to send an API request for the `CBA.AX` ticker over a specified 2-year period. The `yfinance` library conveniently handles the first layer of preprocessing: it *only* returns data for valid trading days, automatically excluding all weekends and public holidays. This clean, 505-row DataFrame, indexed by date, is then passed directly to the feature engineering stage.

**Code Implementation (from `task7_runner.py`):** This code chunk from the main runner script executes the collection of the stock price data.

```
# Download historical stock data from Yahoo Finance
ticker = 'CBA.AX' # Commonwealth Bank of Australia
end_date = datetime.now()
start_date = end_date - timedelta(days=730) # Last 2 years

stock_df = yf.download(ticker, start=start_date, end=end_date, progress=False)
stock_df.reset_index(inplace=True)
stock_df.columns = [col[0] if isinstance(col, tuple) else col for col in stock_df.columns]

print(f"\n[OK] Downloaded {len(stock_df)} days of stock data")
```

#### Code Explanation (by Chunk):

- **Chunk 1 (Configuration):** The `ticker`, `end_date`, and `start_date` variables are defined. `timedelta(days=730)` ensures a consistent 2-year window of data is requested.
- **Chunk 2 (API Call):** `yf.download(...)` is the core function call. It queries the Yahoo Finance API and returns a Pandas DataFrame containing the Open, High, Low, Close, and Volume data. `progress=False` is used to suppress the console output.
- **Chunk 3 (Cleaning):** `stock_df.reset_index(inplace=True)` moves the `Date` from the DataFrame's index into a regular column. The subsequent line,

`stock_df.columns = ...`, is a helper function to flatten any multi-level column headers that `yfinance` might return, ensuring clean column names like `Close` instead of `('Close', '')`.

---

## Chunk 2: News Data Collection (The Sentiment Source)

**Flow:** The second data stream involves collecting the textual news data. As API limits are a major constraint, the project uses a robust web scraping approach defined in `web_scraper.py`. The flow is as follows:

1. The `WebNewsCollector` class is instantiated.
2. The main `collect_news` method is called, which orchestrates calls to multiple sub-methods (e.g., `_scrape_yahoo_finance`, `_scrape_google_finance`, `_fetch_rss_feeds`).
3. Each sub-method uses the `requests` library to fetch HTML/XML and `BeautifulSoup` to parse the content, extracting the `title`, `description`, `date`, and `source`.
4. All collected articles are aggregated into a single list, converted to a DataFrame, and then passed to the cleaning flow (Chunk 3).
5. Note: The main `task7_runner.py` script loads this data from a pre-collected CSV (`task7_data/news_raw/news_raw.csv`), but the `web_scraper.py` file defines the original collection logic.

**Code Implementation (from `web_scraper.py`):** This code chunk shows the scraping logic for a public RSS feed, which is a reliable method for collecting structured news data without an API key.

```
def _fetch_rss_feeds(self, company_name: str, max_articles: int) -> List[Dict]:  
    ....  
    Fetch news from financial RSS feeds  
    ....  
    articles = []  
  
    rss_feeds = [  
        'https://www.ft.com/rss/companies/banks',  
        'https://feeds.bloomberg.com/markets/news.rss',  
        'https://www.reuters.com/rssFeed/businessNews',
```

```
]

for feed_url in rss_feeds:
    try:
        # 1. Fetch the RSS (XML) data
        response = self.session.get(feed_url, timeout=10)
        if response.status_code != 200:
            continue

        # 2. Parse the XML
        soup = BeautifulSoup(response.content, 'xml')
        items = soup.find_all('item')

        for item in items:
            # 3. Extract relevant fields
            title = item.find('title').get_text(strip=True)

            # 4. Filter by relevance
            if company_name.lower() not in title.lower():
                continue

            description = item.find('description').get_text(strip=True) if item.find('description')
            else ""

            link = item.find('link').get_text(strip=True) if item.find('link') else ""
            pub_date_str = item.find('pubDate').get_text(strip=True) if item.find('pubDate') else ""
            date = self._parse_date(pub_date_str) # (Helper function to parse date)

            # 5. Append to list
            articles.append({
                'date': date,
```

```

        'title': title,
        'description': description,
        'content': description,
        'source': 'RSS Feed',
        'url': link
    })
except Exception as e:
    continue
return articles

```

### Code Explanation (by Chunk):

- **Chunk 1 (Iteration):** The code iterates through a predefined list of `rss_feeds`.
  - **Chunk 2 (Fetch & Parse):** It uses `self.session.get` to download the raw XML content of the feed. `BeautifulSoup(response.content, 'xml')` is then used to parse this content, and `soup.find_all('item')` creates a list of all news articles in that feed.
  - **Chunk 3 (Extraction & Filtering):** For each `item`, the code extracts the `title`, `description`, and `pubDate`. A crucial business logic step is performed: `if company_name.lower() not in title.lower(): continue`. This filters out irrelevant articles, ensuring only news related to "Commonwealth Bank" is kept.
  - **Chunk 4 (Aggregation):** The cleaned, relevant data is appended as a dictionary to the `articles` list, which is returned at the end.
- 

### Chunk 3: News Text Cleaning (Sanitization)

**Flow:** The 99 raw articles are unusable for sentiment analysis. They are "dirty" with digital noise. This flow, defined in `news_collector.py`, sanitizes them.

1. **Input:** A raw text string (e.g., a `title` or `description`) is passed to the `_clean_text` function.
2. **Normalization:** The text is converted to lowercase for consistency (e.g., "Profit" and "profit" are treated as the same word).
3. **Sanitization Pipeline:** The text is piped through a series of regular expression (regex) filters to strip noise in a specific order:
  - First, HTML tags (e.g., `<p>`, `<strong>`) are removed.
  - Second, URLs (e.g., `http://...`) are removed.
  - Third, all non-alphanumeric characters (punctuation, symbols) are removed.
4. **Whitespace Collapsing:** All excess whitespace (tabs, newlines, multiple spaces) is

collapsed into single spaces.

5. **Output:** A clean, normalized string containing only words is returned, ready for sentiment analysis.

**Code Implementation (from `news_collector.py`):** The `_clean_text` method implements this sanitization pipeline.

```
def _clean_text(self, text: str) -> str:  
    """  
    Clean individual text field  
    """  
  
    if pd.isna(text) or not isinstance(text, str):  
        return ""  
  
    # --- Chunk 3.1: HTML Tag Removal ---  
    # Use BeautifulSoup to parse and remove any HTML  
    text = BeautifulSoup(text, 'html.parser').get_text()  
  
    # --- Chunk 3.2: URL and Special Char Removal ---  
    # Remove URLs  
    text = re.sub(r'http\S+|www\S+|https\S+', '', text, flags=re.MULTILINE)  
    # Remove email addresses  
    text = re.sub(r'\S+@\S+', '', text)  
    # Remove special characters but keep basic punctuation  
    text = re.sub(r'[^a-zA-Z0-9\s.,!?-]', '', text)  
  
    # --- Chunk 3.3: Whitespace Collapsing ---  
    # Remove extra whitespace, newlines, tabs  
    text = ' '.join(text.split())  
  
    return text.strip()
```

## Code Explanation (by Chunk):

- **Chunk 3.1 (HTML Removal):** `BeautifulSoup(text, 'html.parser').get_text()` is a robust way to strip all HTML tags from the text, leaving only the human-readable content.
  - **Chunk 3.2 (Regex Sanitization):** This chunk is a pipeline of regular expressions. `re.sub(r'http\S+...', '', ...)` finds and removes all URLs. `re.sub(r'[^\w\s.,!?-]', '', text)` is a key filter that removes any character that is *not* a letter, number, whitespace, or basic punctuation, effectively cleaning out noise.
  - **Chunk 3.3 (Whitespace Collapsing):** `text.split()` breaks the string into a list of words, which discards all whitespace. `' '.join(...)` reassembles them with single spaces. This is a highly efficient way to normalize all whitespace.
- 

## Chunk 4: Time Alignment (The Critical Synchronization Step)

**Flow:** This is the most critical preprocessing step, addressing the temporal mismatch between the 24/7 news cycle and the Mon-Fri stock market.

1. **Problem:** A positive news article published on a Sunday must influence the *next* trading day (Monday), not the non-existent "Sunday" stock price.
2. **Solution (Flow):** We must "push" all weekend news forward to the next valid trading day.
3. **Implementation:** The `_align_to_trading_days` function (from `news_collector.py`) manages this:
  - **Detection:** It identifies the day of the week (Monday=0, Sunday=6) for each article.
  - **Offset Logic:** It applies business logic to calculate an offset. If `weekday == 5` (Saturday), the offset is +2 days. If `weekday == 6` (Sunday), the offset is +1 day.
  - **Synchronization:** It applies this offset, shifting all weekend dates to the following Monday.
4. **Output:** A DataFrame of news articles where all dates are now aligned with the trading calendar.

**Code Implementation (from `news_collector.py`):** This function is essential for preventing look-ahead bias.

```
def _align_to_trading_days(self, df: pd.DataFrame) -> pd.DataFrame:
```

```
    ....
```

```

Align news dates to trading days (Move weekends to next Monday)

print(" [6] Aligning to trading days...")

df_aligned = df.copy()

# --- Chunk 4.1: Detection ---
# Get the day of the week (Monday=0, Sunday=6)
df_aligned['weekday'] = df_aligned['date'].dt.dayofweek

# --- Chunk 4.2: Offset Logic ---
# Create an offset column, defaulting to 0
df_aligned['days_to_add'] = 0
# Find all Saturdays (5) and set their offset to +2 days
df_aligned.loc[df_aligned['weekday'] == 5, 'days_to_add'] = 2
# Find all Sundays (6) and set their offset to +1 day
df_aligned.loc[df_aligned['weekday'] == 6, 'days_to_add'] = 1

# --- Chunk 4.3: Synchronization ---
# Apply the offset to the date to get the correct trading date
df_aligned['date'] = df_aligned['date'] + pd.to_timedelta(df_aligned['days_to_add'],
unit='D')

# --- Chunk 4.4: Cleanup ---
# Remove the temporary helper columns
df_aligned = df_aligned.drop(columns=['weekday', 'days_to_add'])

return df_aligned

```

#### Code Explanation (by Chunk):

- **Chunk 4.1 (Detection):** The flow starts by using the Pandas `.dt.dayofweek` accessor to get the weekday number for each article's date.

- **Chunk 4.2 (Offset Logic):** This is the core business rule. It uses `.loc` to find all rows where `weekday == 5` (Saturday) and assigns `2` to their `days_to_add` column. It does the same for `weekday == 6` (Sunday), assigning `1`. All other days (Mon-Fri) remain `0`.
  - **Chunk 4.3 (Synchronization):** `pd.to_timedelta` is used to safely add the `days_to_add` offset to the original `date`. This operation correctly pushes all Saturday and Sunday dates forward to the following Monday, ensuring the data is temporally sound.
  - **Chunk 4.4 (Cleanup):** The helper columns are dropped, leaving the DataFrame clean.
- 

## Chunk 5: Data Synchronization & Missing Data Handling

**Flow:** This is the final preprocessing flow where the two streams (Stock prices and News sentiment) are merged into a single dataset. This logic is handled by the `SentimentFeatureBuilder` class in `feature_builder.py`.

1. **Inputs:**
  - `stock_with_tech`: The 505-day DataFrame of stock prices (from Chunk 1).
  - `sentiment_df`: The DataFrame of news, after Section III's sentiment analysis and daily aggregation (e.g., ~50-60 days that have news).
2. **Merge:** The two DataFrames are merged using a `left` join, anchored to the `stock_with_tech` DataFrame.
3. **Problem:** This `left` join keeps all 505 trading days but creates `NaN` (Missing) values in the sentiment columns for the 450+ days that had *no* news.
4. **Business Assumption:** The flow makes a critical assumption: **No News = Neutral Sentiment**.
5. **Implementation:** This assumption is implemented by filling all `NaN` values in sentiment columns with `0`.
6. **Output:** A single, synchronized, `NaN`-free DataFrame of 505 rows, ready for modeling.

**Code Implementation (from `feature_builder.py`):** This code snippet shows the merge and the handling of missing data.

```
def merge_sentiment(self, sentiment_df: pd.DataFrame) -> pd.DataFrame:

    # (Stock data is loaded and technical features added first)
    stock_with_tech = self.add_technical_indicators(self.stock_df.copy())

    # Ensure date columns are matching types
    stock_with_tech['date'] = pd.to_datetime(stock_with_tech['date']).dt.date
```

```

sentiment_df['date'] = pd.to_datetime(sentiment_df['date']).dt.date

# --- Chunk 5.1: The Merge ---
# Merge on date, using 'left' to keep all trading days
merged = pd.merge(
    stock_with_tech,
    sentiment_df,
    on='date',
    how='left' # <-- Keeps all 505 stock days
)

# --- Chunk 5.2: Missing Data Handling ---
# Apply the "No News = Neutral" assumption
sentiment_cols = [col for col in sentiment_df.columns if col != 'date']

# Fill sentiment_score with 0 (neutral)
if 'sentiment_score' in merged.columns:
    merged['sentiment_score'] = merged['sentiment_score'].fillna(0)

# Fill counts with 0
if 'article_count' in merged.columns:
    merged['article_count'] = merged['article_count'].fillna(0)

# Fill volatility metrics with 0
if 'sentiment_std' in merged.columns:
    merged['sentiment_std'] = merged['sentiment_std'].fillna(0)

return merged

```

#### Code Explanation (by Chunk):

- **Chunk 5.1 (The Merge):** `pd.merge(..., how='left')` is the key. It anchors the

merge to the `stock_with_tech` DataFrame (the 505 trading days) and only attaches sentiment data from `sentiment_df` where the dates match. All other rows in `stock_with_tech` are kept, but their sentiment columns are filled with `NaN`.

- **Chunk 5.2 (Missing Data Handling):** This chunk implements the "No News = Neutral" assumption. The `.fillna(0)` calls are critical: they replace all `NaN` values (days with no news) with `0`. This correctly informs the model that on those specific days, the sentiment signal was neutral (0.0), not positive or negative.

## Section III. Sentiment Analysis

This section details the methodology used to convert the cleaned textual data (from Section II) into quantitative sentiment features. This process involves two main stages: (1) calculating sentiment for individual articles, and (2) aggregating these scores to a daily level to align with the stock price data.

### Chunk 1: Tool Selection and Rationale

**Flow:** The first step in the sentiment analysis flow is selecting the right tool. The project requirements ([Tasks C.7 - Extension.pdf](#)) allow for experimenting with different tools. A key decision was made to *not* use complex, deep learning models like FinBERT (which is reserved for Section VI: Independent Research) in the main pipeline.

**Approach & Rationale:** For the primary pipeline (`task7_runner.py`), a **lexicon-based (dictionary-based)** model was chosen.

1. **Tool:** The project uses `TextBlob` (called via the `SentimentAnalyzer` class, as seen in `task7_runner.py`).
2. **Why TextBlob?**
  - **Speed & Efficiency:** It is extremely fast, processing 99 articles in milliseconds. This is crucial for a rapid pipeline.
  - **No GPU Required:** Unlike FinBERT, it runs on any CPU.
  - **Simplicity:** It provides a single, standardized polarity score from -1.0 (highly negative) to +1.0 (highly positive), which is ideal for feature engineering.
3. **Instantiation:** The main runner script instantiates the analyzer, explicitly choosing the '`lexicon`' method.

**Code Implementation (from `task7_runner.py`):** This code shows the instantiation of the chosen sentiment analysis tool.

```
# Task 7 custom modules for sentiment-based prediction  
# ...  
from task7_sentiment.sentiment_analyzer import SentimentAnalyzer
```

```

# ...

def main():
    # ... (Stage 2) ...
    print("\n" + "*80)
    print("STAGE 2: SENTIMENT ANALYSIS")
    print("*80)

# --- Chunk 1.1: Tool Instantiation ---
# Initialize sentiment analyzer with lexicon-based method (fast and reliable)
analyzer = SentimentAnalyzer(primary_model='lexicon')

```

- **Code Explanation:**

- `from task7_sentiment.sentiment_analyzer import SentimentAnalyzer`: This line imports the custom class responsible for handling sentiment analysis.
  - `analyzer = SentimentAnalyzer(primary_model='lexicon')`: This line creates an instance of the analyzer. By passing `primary_model='lexicon'`, we are configuring the system to use a fast, dictionary-based approach (which in turn calls `textblob`) rather than a heavy neural network.
- 

## Chunk 2: Article-Level Sentiment Calculation Flow

**Flow:** Once the tool is selected, the flow proceeds to analyze each of the 99 news articles individually.

1. **Input:** The DataFrame `news_df` containing the 99 cleaned articles (from Section II).
2. **Iteration:** The flow uses the efficient Pandas `.apply()` method to iterate over the `title` column of the DataFrame.
3. **Analysis:** For each title, the `.apply()` method passes the text (as `x`) to the `analyzer.analyze` function.
4. **Scoring:** The `analyzer` (using TextBlob) processes the text and returns a single polarity score (e.g., `0.35`).
5. **Output:** A new column, `sentiment_score`, is created in the DataFrame, storing the score for each individual article.

**Code Implementation (from `task7_runner.py`):** This chunk of code performs the sentiment analysis on all 99 articles.

```
# --- Chunk 2.1: Analyze each article title ---  
# Apply the sentiment analyzer to the 'title' column  
# The lambda function calls the analyzer for each row  
news_df['sentiment_score'] = news_df['title'].apply(  
    lambda x: analyzer.analyze(str(x), method='textblob')  
)  
  
# --- Chunk 2.2: Categorize for interpretability ---  
# (This step is for logging/stats, not for feature engineering)  
news_df['sentiment_category'] = news_df['sentiment_score'].apply(  
    lambda x: 'positive' if x > 0.05 else ('negative' if x < -0.05 else 'neutral')  
)
```

#### Code Explanation (by Chunk):

- **Chunk 2.1 (Analysis):** `news_df['title'].apply(...)` is the core of this flow.
  - `lambda x::`: This creates a small, anonymous function that runs for each title (`x`).
  - `analyzer.analyze(str(x), method='textblob')`: This is the call to the analysis engine. It takes the title string, runs TextBlob's polarity algorithm on it, and returns the float score. The result is saved in the `sentiment_score` column.
- **Chunk 2.2 (Categorization):** This second `.apply()` is used for verification and reporting. It converts the continuous score (e.g., `0.35`) into a simple category ("positive"), which is then used to print the distribution (e.g., "Positive: 65%, Negative: 7%").

---

#### Chunk 3: Daily Aggregation Flow (Critical Requirement)

**Flow:** This is the most important flow in Section III, as it addresses a key requirement from [Tasks C.7 – Extension.pdf](#): "scores should be aggregated at a daily level."

1. **Problem:** On a given trading day (e.g., 2024-05-10), there might be 5 news articles with 5 different sentiment scores. The model, however, needs one set of sentiment features for that single day.

2. **Solution (Flow):** We aggregate all articles from the *same day* into a single row.
3. **Implementation:** The flow uses the Pandas `groupby()` function on the (already time-aligned) `date` column.
4. **Aggregation:** The `.agg()` function is then called to compute multiple statistics for all articles within that group (day):
  - o `mean`: The average sentiment score (e.g., `sentiment_mean`). This is the primary feature.
  - o `std`: The standard deviation of scores (e.g., `sentiment_std`). This measures sentiment *volatility* or *disagreement* (a high `std` means some news was very positive and some very negative).
  - o `count`: The number of articles (`article_count`).
  - o `positive_ratio`: The percentage of positive articles.
5. **Output:** A new DataFrame, `daily_sentiment`, where each row represents a single, unique trading day with its computed sentiment features.

**Code Implementation (from `task7_runner.py`):** This code block aggregates the article-level scores into daily features.

```
# --- Chunk 3.1: Ensure date format ---

# Aggregate sentiment by day
news_df['date'] = pd.to_datetime(news_df['date'])

# --- Chunk 3.2: Group and Aggregate ---
daily_sentiment = news_df.groupby(news_df['date'].dt.date).agg({
    'sentiment_score': ['mean', 'std', 'count'], # Mean, volatility, article count
    'sentiment_category': lambda x: (x == 'positive').sum() / len(x) if len(x) > 0 else 0 # Positive ratio
}).reset_index()
```

**Code Explanation (by Chunk):**

- **Chunk 3.1 (Date Format):** Ensures the `date` column is a `datetime` object, which is required for the `.groupby()` operation.
- **Chunk 3.2 (Group & Agg):** This is the core logic.
  - o `news_df.groupby(news_df['date'].dt.date)`: This groups all 99 articles by their unique date.
  - o `.agg({ ... })`: This dictionary tells Pandas what to compute for each group.
  - o `'sentiment_score': ['mean', 'std', 'count']`: This is the key instruction. It tells Pandas to take the `sentiment_score` column for that day's articles and calculate its mean, standard deviation, and count.

- `lambda x: ...`: This calculates the ratio of positive articles for that day.
- 

## Chunk 4: Sentiment Feature Creation

**Flow:** The output from Chunk 3 (`daily_sentiment`) has messy, multi-level column names (e.g., `('sentiment_score', 'mean')`). The flow must clean this up to create usable feature columns.

1. **Input:** The `daily_sentiment` DataFrame with multi-level columns.
2. **Flattening:** The column names are flattened into simple, single-level names.
3. **Feature Selection:** As defined in `task7_runner.py`, only the most important aggregated features are selected for the model: `sentiment_mean` and `sentiment_std`.
4. **Output:** A clean DataFrame with columns `date`, `sentiment_mean`, `sentiment_std`, etc., ready to be merged (in Chunk 5) with the stock data.

**Code Implementation (from `task7_runner.py`):** This code flattens the column names and identifies the final features.

```
# --- Chunk 4.1: Flatten multi-level columns ---
daily_sentiment.columns = ['date', 'sentiment_mean', 'sentiment_std',
                           'article_count', 'positive_ratio']

# --- Chunk 4.2: Handle potential NaNs ---
# Fill std with 0 for single-article days (std is NaN for 1 value)
daily_sentiment['sentiment_std'] = daily_sentiment['sentiment_std'].fillna(0)

# ... (Later in Stage 3) ...

# --- Chunk 4.3: Define Sentiment Feature List ---
# Identify feature columns by prefix
sentiment_features = [c for c in full_df.columns if c.startswith('sentiment_')]
# (This list will resolve to ['sentiment_mean', 'sentiment_std', ...])
```

## Code Explanation (by Chunk):

- **Chunk 4.1 (Flattening):** `daily_sentiment.columns = [...]` is a direct way to

rename the complex ('sentiment\_score', 'mean') headers to the simple `sentiment_mean`, which is required for the model.

- **Chunk 4.2 (Handle NaNs):** This is a subtle but important data cleaning step. If a day had only one article, its standard deviation (`std`) is `NaN` (mathematically undefined). `fillna(0)` correctly sets the sentiment volatility to 0 for those days, as there was no disagreement.
  - **Chunk 4.3 (Feature List):** The line `sentiment_features = [c for c in ...]` programmatically selects all columns that start with `sentiment_`, resulting in the final list of features (`sentiment_mean`, `sentiment_std`, etc.) that will be fed to the model.
- 

## Chunk 5: Handling News-less Days (Merge & FillNA)

**Flow:** This final flow connects the sentiment data (now aggregated by day) to the stock data backbone.

1. **Problem:** The `stock_df` has 505 trading days. The `daily_sentiment` DataFrame only has ~50-60 rows (days with news). We must create sentiment features for the 450+ "news-less" days.
2. **Merge:** The flow, executed in `task7_runner.py` (Stage 3), performs a `left` merge, anchoring on the 505-day `stock_df`.
3. **Result:** This merge creates `NaN` (Missing) values in `sentiment_mean` and `sentiment_std` for all days that had no news.
4. **Business Assumption:** A critical assumption is made: **No News = Neutral Sentiment**.
5. **Implementation:** The flow uses `.fillna(0)` to replace all `NaN` values in the sentiment columns with `0`.
6. **Output:** A single, 505-row DataFrame where every trading day has a valid sentiment score (either the calculated score or 0 for neutral).

**Code Implementation (from `task7_runner.py`):** This code (from Stage 3 of the runner) shows the merge and the handling of missing data.

```
# --- Chunk 5.1: The Merge ---  
  
# Merge stock data with sentiment data on date  
  
# Left join ensures we keep all stock days (even without news)  
stock_df['date'] = pd.to_datetime(stock_df['Date']).dt.date  
daily_sentiment['date'] = pd.to_datetime(daily_sentiment['date']).dt.date  
  
full_df = stock_df.merge(daily_sentiment, on='date', how='left')
```

```
# --- Chunk 5.2: Missing Data Handling ---  
  
# Fill missing sentiment values (days without news) with neutral values  
# This is important: missing news doesn't mean negative sentiment!  
full_df['sentiment_mean'] = full_df['sentiment_mean'].fillna(0) # Neutral sentiment  
full_df['sentiment_std'] = full_df['sentiment_std'].fillna(0) # No volatility  
full_df['article_count'] = full_df['article_count'].fillna(0) # No articles  
full_df['positive_ratio'] = full_df['positive_ratio'].fillna(0.5) # 50% positive (neutral)
```

### Code Explanation (by Chunk):

- **Chunk 5.1 (The Merge):** `how='left'` is the most important parameter. It preserves all 505 rows from `stock_df` (the left table) and only attaches data from `daily_sentiment` (the right table) where the dates match.
- **Chunk 5.2 (Missing Data Handling):** This chunk implements the "No News = Neutral" assumption. `full_df[ 'sentiment_mean' ].fillna(0)` replaces all `NaN` values with `0`. This is a crucial step that ensures the model has a complete dataset and correctly interprets days without news as having a neutral (0.0) sentiment signal.

## IV. Feature Engineering & Modelling

This section covers the creation of the final features, the definition of the classification target, the preprocessing pipeline, and the comprehensive model training flow.

### Chunk 1: Technical Feature Engineering

**Flow:** The feature engineering flow begins in Stage 3 of `task7_runner.py`. After loading the 505-day `stock_df`, the script *manually* calculates 14 different technical indicators. This is done to create a rich set of features that capture price momentum, trend, and volatility. These calculations are performed directly using Pandas' rolling and math functions.

**Code Implementation (from `task7_runner.py`):** This code snippet shows the direct, manual calculation of these technical features.

```
# (Quoted from: task7_runner.py, STAGE 3)

# --- 1. RETURNS: Measure price changes ---
stock_df['return_1d'] = stock_df['Close'].pct_change()
stock_df['return_5d'] = stock_df['Close'].pct_change(5)

# --- 2. VOLATILITY: Measure price instability ---
stock_df['volatility_20d'] = stock_df['return_1d'].rolling(20).std()

# --- 3. MOVING AVERAGES: Smooth price trends ---
stock_df['ma_20'] = stock_df['Close'].rolling(20).mean()
stock_df['ma_50'] = stock_df['Close'].rolling(50).mean()

# --- 4. RSI (Relative Strength Index): Momentum oscillator ---
delta = stock_df['Close'].diff()
gain = (delta.where(delta > 0, 0)).rolling(14).mean()
loss = (-delta.where(delta < 0, 0)).rolling(14).mean()
rs = gain / loss
stock_df['rsi'] = 100 - (100 / (1 + rs))
```

```

# --- 5. MACD (Moving Average Convergence Divergence) ---
exp1 = stock_df['Close'].ewm(span=12, adjust=False).mean()
exp2 = stock_df['Close'].ewm(span=26, adjust=False).mean()
stock_df['macd'] = exp1 - exp2
stock_df['macd_signal'] = stock_df['macd'].ewm(span=9, adjust=False).mean()

```

### Code Explanation:

- **Chunks 1-3:** These use standard Pandas functions. `.pct_change(5)` calculates the 5-day return, and `.rolling(20).std()` calculates the 20-day rolling volatility.
  - **Chunk 4 (RSI):** This implements the standard RSI formula by calculating the average `gain` and `loss` over a 14-day window.
  - **Chunk 5 (MACD):** This calculates the 12-day and 26-day Exponential Moving Averages (EMA) and finds their difference (`macd`) and the 9-day EMA of that difference (`macd_signal`).
- 

### Chunk 2: Target Variable Creation (The Core Task)

**Flow:** This is the most important step for Task C.7. Instead of predicting the *price* (a regression problem), we are predicting the *direction* (a classification problem). The flow is simple but critical:

1. Take the `Close` price column.
2. Use `.shift(-1)` to "look" one day into the future and pull that price onto the current day's row.
3. Compare that future price to the current day's price.
4. If `Future_Close > Current_Close`, set the target to `1` (UP). Otherwise, set it to `0` (DOWN).
5. This process creates `Nans` (e.g., on the last row), which are then dropped.

**Code Implementation (from `task7_runner.py`):** This single line of code transforms the project from regression to classification.

```
# (Quoted from: task7_runner.py, STAGE 3)
```

```

# --- Create binary target variable: 1 if price goes UP tomorrow, 0 if DOWN ---
# shift(-1) looks at next day's price (future prediction target)
full_df['target'] = (full_df['Close'].shift(-1) > full_df['Close']).astype(int)

```

```
# Remove rows with NaN (caused by rolling windows and shift operations)
full_df = full_df.dropna()
```

### Code Explanation:

- `full_df['Close'].shift(-1)`: This is the key. It pulls tomorrow's closing price onto today's row.
  - `(...) > full_df['Close']`: This boolean comparison returns `True` (if the price went up) or `False`.
  - `.astype(int)`: This converts `True` to `1` and `False` to `0`, creating our binary target.
  - `full_df = full_df.dropna()`: This cleans the DataFrame, removing the first 50 rows (due to the 50-day MA) and the very last row (due to the `.shift(-1)`).
- 

### Chunk 3: Preprocessing Pipeline (SMOTE & Scaling)

**Flow:** Before training, the data (now 456 samples) must be preprocessed. This flow is defined in the `SentimentClassifierTrainer` class within `classifier_models.py`.

1. **Temporal Split:** First, the data is split into training (80%) and testing (20%) sets *temporally* (no shuffling) in `task7_runner.py`.
2. **Scaling:** The `preprocess_data` function is called. It fits a `StandardScaler` to the *training data only* to learn its mean/std. It then uses this *fitted* scaler to transform both the train and test sets.
3. **Balancing (SMOTE):** The training set is imbalanced (149 DOWN vs 215 UP). To fix this, `SMOTE` (Synthetic Minority Over-sampling Technique) is applied *only to the scaled training set*. It synthetically creates new "DOWN" samples until the classes are balanced (215 UP vs 215 DOWN).
4. **Output:** The flow returns the `(X_train_scaled_balanced, y_train_balanced, X_test_scaled)`.

**Code Implementation (from `classifier_models.py`):** This function shows the complete, correct preprocessing pipeline.

```
# (Quoted from: classifier_models.py)

def preprocess_data(self, X_train: pd.DataFrame, y_train: pd.Series,
                   X_test: pd.DataFrame = None,
                   experiment_name: str = 'default') -> Tuple:
```

```

# --- Chunk 3.1: Scaling ---
print(" [1] Scaling features...")
scaler = StandardScaler()

# Fit on TRAIN, transform TRAIN
X_train_scaled = scaler.fit_transform(X_train)
# Transform TEST using TRAIN's scaler
X_test_scaled = scaler.transform(X_test) if X_test is not None else None

self.scalers[experiment_name] = scaler # Store the scaler

# --- Chunk 3.2: SMOTE (on training data ONLY) ---
if self.use_smote:
    print(" [2] Applying SMOTE for class balance...")
    unique, counts = np.unique(y_train, return_counts=True)
    print(f" Before SMOTE: {dict(zip(unique, counts))}")

    smote = SMOTE(random_state=self.random_state)
    X_train_resampled, y_train_resampled = smote.fit_resample(
        X_train_scaled, y_train
    )

    unique, counts = np.unique(y_train_resampled, return_counts=True)
    print(f" After SMOTE: {dict(zip(unique, counts))}")

    X_train_processed = X_train_resampled
    y_train_processed = y_train_resampled
    # ... (else statement)
    return X_train_processed, y_train_processed, X_test_scaled

```

## Code Explanation:

- `scaler.fit_transform(X_train)`: Correctly fits *and* transforms the training data.
  - `scaler.transform(X_test)`: Correctly applies the *same* transformation (using the training data's mean/std) to the test data. This prevents data leakage.
  - `smote.fit_resample(X_train_scaled, y_train)`: This is the crucial SMOTE step. It is applied *after* scaling and *only* on the training set.
- 

## Chunk 4: The 3x7 Experimental Design (Training Flow)

**Flow:** This project runs a comprehensive 3x7 factorial experiment (21 models total), orchestrated by two scripts: `task7_runner.py` and `task7_extended_models.py`.

1. **`task7_runner.py (Stage 4)`:** This script trains the first 9 models. It loops through 3 feature sets (Technical, Sentiment, Combined) and 3 algorithms (Logistic, RF, XGBoost).
2. **`task7_extended_models.py`:** This script is then run. It loads the *same* preprocessed data and trains the remaining 12 models: 4 new algorithms (SVM-Linear, SVM-RBF, Gradient Boosting, MLP) on the same 3 feature sets.
3. **`run_all.py`:** This script simply calls the two runner scripts above in sequence to execute the full 21-model experiment.
4. **Output:** All 21 models are saved as `.pkl` files in the `task7_models/` directory, and their results are merged into `evaluation_metrics.json`.

**Code Implementation (from `task7_runner.py & task7_extended_models.py`):** First, the 3x3 loop in `task7_runner.py` defines the 3 feature sets and trains the 9 base models.

```
# (Quoted from: task7_runner.py, STAGE 4)

# Define three feature sets for comparison
feature_sets = {
    'technical_only': technical_features,
    'sentiment_only': sentiment_features,
    'combined': technical_features + sentiment_features
}
```

```

# Three classification algorithms
algorithms = ['logistic', 'random_forest', 'xgboost']

trainer = SentimentClassifierTrainer(use_smote=True, random_state=42)

# Loop through each feature set
for feat_name, feat_cols in feature_sets.items():
    # ... (Get X_train, X_test, y_train, y_test for this set) ...
    X_train_proc, y_train_proc, X_test_proc = trainer.preprocess_data(...)

    # Loop through each algorithm
    for algo in algorithms:
        model_name = f"{feat_name}_{algo}"
        model = trainer.train_model(X_train_proc, y_train_proc, model_type=algo, ...)
        # ... (Save model and store predictions) ...

```

Second, `task7_extended_models.py` calls its internal functions to train the other 12 models.

```

# (Quoted from: task7_extended_models.py, main())

# 1. Train SVM models (6 models: 2 kernels × 3 feature sets)
svm_results, svm_models = train_svm_models(full_df, tech_feats, sent_feats, y,
split_idx)
all_new_results.extend(svm_results)

# 2. Train Gradient Boosting models (3 models: 1 algorithm × 3 feature sets)
gb_results, gb_models = train_gradientboosting_models(full_df, tech_feats,
sent_feats, y, split_idx)
all_new_results.extend(gb_results)

# 3. Train MLP Neural Network models (3 models: 1 architecture × 3 feature sets)
mlp_results, mlp_models = train_mlp_models(full_df, tech_feats, sent_feats, y,

```

```

    split_idx)

    all_new_results.extend(mlp_results)

    # Merge with existing results (9 from main + 12 from extended = 21 total)
    update_evaluation_results(all_new_results)

```

### Code Explanation:

- The nested `for` loops in `task7_runner.py` clearly show the 3x3 experimental design.
- `task7_extended_models.py` systematically calls its `train_...` functions to train the advanced models on the *exact same* 3 feature sets.
- `update_evaluation_results` (in `task7_extended_models.py`) merges the JSON results, creating one master file with all 21 models for comparison.

## V. Evaluations

This section concisely covers how the 21 trained models were evaluated and, most importantly, how the value of sentiment was assessed.

### Chunk 1: Comprehensive Metrics Calculation

**Flow:** The evaluation flow (Stage 5 in `task7_runner.py`) iterates through every prediction made by the 9 base models. The `task7_extended_models.py` script does the same for its 12 models.

1. For each model, it retrieves the true labels (`y_test`) and the model's predictions (`y_pred`).
2. It passes these arrays to a suite of functions from `sklearn.metrics` (e.g., `accuracy_score`, `precision_score`, `recall_score`, `f1_score`).
3. These metrics are stored in a dictionary.
4. Finally, `update_evaluation_results` (in `task7_extended_models.py`) merges all 21 dictionaries and saves them as `evaluation_metrics.json` and `model_comparison.csv`.

**Code Implementation (from `task7_runner.py`):** This code from Stage 5 shows the comprehensive metrics calculation.

```
# (Quoted from: task7_runner.py, STAGE 5)
```

```
all_results = [] # Store all evaluation results
```

```

# Evaluate each model on test set
for model_name, preds in all_predictions.items():
    y_test = preds['y_test'] # True labels
    y_pred = preds['y_pred'] # Predicted labels

    # Calculate all required metrics for Task C.7
    metrics = {
        'model': model_name,
        'feature_set': preds['feature_set'],
        'algorithm': preds['algorithm'],
        'accuracy': accuracy_score(y_test, y_pred),
        'precision': precision_score(y_test, y_pred, zero_division=0),
        'recall': recall_score(y_test, y_pred, zero_division=0),
        'f1': f1_score(y_test, y_pred, zero_division=0)
    }

    # Confusion matrix: [[TN, FP], [FN, TP]]
    cm = confusion_matrix(y_test, y_pred)
    metrics['confusion_matrix'] = cm.tolist()

    all_results.append(metrics)

```

#### Code Explanation:

- This block systematically calculates the four key classification metrics (Accuracy, Precision, Recall, F1).
  - **Precision ( $TP / (TP+FP)$ )** answers: "Of all the days we predicted UP, what percentage was actually UP?"
  - **Recall ( $TP / (TP+FN)$ )** answers: "Of all the days that actually went UP, what percentage did we catch?"
  - **F1 Score** is the harmonic mean of Precision and Recall, providing a single, balanced score for comparison.
-

## Chunk 2: The Baseline Comparison (The "So What?")

**Flow:** This is the most important evaluation step, as it directly answers the task's central question: "does sentiment add value?" This flow is executed at the end of `task7_runner.py`.

1. The code filters the `all_results` list to find the best model (by F1 score) from the `technical_only` group. This is the **Baseline**.
2. It does the same to find the best model from the `sentiment_only` and `combined` groups.
3. It then prints a direct comparison of these three "champion" models.
4. Finally, it calculates the percentage improvement of the best sentiment-based model over the technical baseline, providing a clear, quantitative answer. Note: `task7_advanced_evaluation.py` also runs this logic in its `create_summary_report`.

**Code Implementation (from `task7_runner.py`):** This block provides the final conclusion for the report.

```
# (Quoted from: task7_runner.py, BASELINE COMPARISON)

# Best technical-only model (BASELINE - no sentiment features)
tech_results = [r for r in all_results if r['feature_set'] == 'technical_only']
best_tech = max(tech_results, key=lambda x: x['f1']) if tech_results else None

# Best sentiment-only model
sent_results = [r for r in all_results if r['feature_set'] == 'sentiment_only']
best_sent = max(sent_results, key=lambda x: x['f1']) if sent_results else None

# ... (Find best_comb) ...

# Calculate percentage improvement (sentiment vs baseline)
if best_tech and best_sent:
    # ... (Print F1 scores) ...
    imp_f1 = (best_sent['f1'] - best_tech['f1']) / best_tech['f1'] * 100
    print(f"F1 Score Improvement: {imp_f1:+.1f}%")
```

```
if best_sent['f1'] > best_tech['f1']:
    print("\n✓ CONCLUSION: Sentiment features ADD SIGNIFICANT VALUE!")
```

### Code Explanation:

- `best_tech = max(..., key=lambda x: x['f1'])`: This line efficiently finds the dictionary in the `tech_results` list that has the highest F1 score. This is our baseline to beat.
  - `imp_f1 = ...`: This line calculates the percentage change, directly quantifying the value added by sentiment.
  - The final `if` statement provides the project's definitive conclusion.
- 

### Chunk 3: Visualization & Reporting Flow

**Flow:** The final step in the evaluation flow is orchestrated by `task7_advanced_evaluation.py`. This script is designed to be run *after* all 21 models have been trained.

1. **Load:** It loads the master `evaluation_metrics.json` (containing all 21 results) and all 21 `.pkl` model files.
2. **Generate Plots:** It calls its internal plotting functions to generate the three key visualizations:
  - `plot_confusion_matrices`: Creates a 2x3 grid of heatmaps for the Top 6 models.
  - `plot_model_comparison`: Creates a 4-in-1 dashboard comparing all 21 models, their feature sets, and algorithms.
  - `plot_feature_importance`: Extracts and plots the feature importances from all tree-based models (RF, XGBoost, Gradient Boosting).
3. **Generate Reports:** It generates two `.txt` files:
  - `classification_reports.txt`: A detailed, multi-page report with metrics for all 21 models.
  - `executive_summary.txt`: A high-level summary of the findings (as seen in `01_executive_summary.md`).
4. **Orchestration:** The `run_all.py` script simply calls this file at the very end to complete the pipeline.

**Code Implementation (from `task7_advanced_evaluation.py`):** The `main()` function of this script shows the entire visualization and reporting flow.

```
# (Quoted from: task7_advanced_evaluation.py)

def main():
    # ...
    print("\n[LOAD] Loading evaluation results...")
    with open('task7_results/evaluation_metrics.json', 'r') as f:
        all_results = json.load(f)

    print("\n[LOAD] Loading trained models...")
    # ... (Loop to load all .pkl files into 'models' dict) ...

    # --- Chunk 3.1: Generate Visualizations ---
    print("\n" + "="*80)
    print("GENERATING VISUALIZATIONS")
    print("="*80)

    plot_confusion_matrices(all_results)
    plot_model_comparison(all_results)

    if models:
        plot_feature_importance(models,
            ['return_1d', ..., 'sentiment_std']) # Pass all 16 feature names

    # --- Chunk 3.2: Generate Text Reports ---
    print("\n" + "="*80)
    print("GENERATING REPORTS")
    print("="*80)

    generate_classification_reports(all_results)
    create_summary_report(all_results)
```

```

# ... (Print summary of outputs) ...

if __name__ == '__main__':
    main()

```

### Code Explanation:

- **Chunk 3.1:** This block calls the three key plotting functions. `plot_confusion_matrices(all_results)` takes the full results, finds the top 6, and plots them. `plot_feature_importance` is passed the `models` themselves (which were loaded from the `.pk1` files) so it can access their `.feature_importances_` attribute.
- **Chunk 3.2:** This block calls the two report-writing functions, which save the final text-based analysis to the `task7_results/` directory, completing the evaluation.

## Section VI. Independent Research Component

To fulfill the independent research requirement, this project went beyond the baseline `TextBlob` model to explore advanced, domain-specific sentiment techniques. These methods are defined in `finbert_tuner.py` and, while not part of the final 21-model experiment, they provide a clear path for future performance improvements.

### Chunk 1: Financial Lexicon Enhancement

**Flow:** The first enhancement was to create a `FinancialLexiconEnhancer` (in `finbert_tuner.py`). This tool *augments* a base sentiment score (like `TextBlob`'s) by scanning the text for domain-specific keywords relevant to Australian banking.

The flow is:

1. A default lexicon is built with terms like 'profit upgrade' (+0.9) or 'APRA investigation' (-0.7).
2. The text is scanned for these terms to create a `lexicon_score`.
3. This score is combined with the `base_sentiment` using a weighted average.

**Code Implementation (from `finbert_tuner.py`):** This function shows the weighted average logic, giving 70% weight to the base model and 30% to the lexicon adjustment.

```
# (Quoted from: finbert_tuner.py)
```

```
def enhance_sentiment(self, text: str, base_sentiment: float) -> Tuple[float, Dict]:
```

```

# ... (Finds matching terms and calculates lexicon_score) ...

if matched_scores:
    lexicon_score = np.mean(matched_scores)
else:
    lexicon_score = 0.0

# --- Chunk 1.1: Weighted Average ---
# Combine scores (weighted average)
# Give more weight to base sentiment (70%), lexicon adds adjustment (30%)
enhanced_sentiment = 0.7 * base_sentiment + 0.3 * lexicon_score

# Ensure within bounds [-1, 1]
enhanced_sentiment = np.clip(enhanced_sentiment, -1.0, 1.0)

# ... (Return enhanced_sentiment and details) ...

```

### Code Explanation:

- **Chunk 1.1 (Weighted Average):** This logic (`0.7 * base_sentiment + 0.3 * lexicon_score`) provides a simple way to "nudge" the generic sentiment score in the correct direction based on highly specific financial terms, making the model more domain-aware.
- 

### Chunk 2: FinBERT Fine-Tuning Strategy

**Flow:** The second and most advanced technique is the `FinBERTFineTuner` class. The flow is designed to create a "specialist" model:

1. **Load:** It loads the powerful, pre-trained `ProsusAI/finbert` model.
2. **Justification:** This base model is trained on *general* financial news (mostly US/EU). To improve its accuracy for *our* specific task, we must fine-tune it on *our* data (CBA/Australian news).
3. **Auto-Label:** Since we have no human-labeled data, the `prepare_training_data` function uses the pre-trained FinBERT model to create initial "weak" labels for our 99

articles.

4. **Fine-Tune:** The `fine_tune` function then re-trains the model on this small, domain-specific dataset, making it an expert on Australian banking terminology.

**Code Implementation (from `finbert_tuner.py`):** The `fine_tune` method uses the Hugging Face `Trainer` to re-train the model.

```
# (Quoted from: finbert_tuner.py)
```

```
def fine_tune(self, train_df: pd.DataFrame, val_df: pd.DataFrame = None,
             epochs: int = 3, batch_size: int = 8,
             learning_rate: float = 2e-5):

    # --- Chunk 2.1: Create Datasets ---
    train_dataset = self.create_dataset(train_df)
    val_dataset = self.create_dataset(val_df) if val_df is not None else None

    # --- Chunk 2.2: Define Training Arguments ---
    training_args = TrainingArguments(
        output_dir=self.output_dir,
        num_train_epochs=epochs,
        per_device_train_batch_size=batch_size,
        learning_rate=learning_rate,
        evaluation_strategy='epoch' if val_dataset else 'no',
        # ... (Other parameters) ...
    )

    # --- Chunk 2.3: Initialize Trainer ---
    trainer = Trainer(
        model=self.model,
        args=training_args,
        train_dataset=train_dataset,
        eval_dataset=val_dataset,
```

```

# ... (Other parameters) ...
)

# --- Chunk 2.4: Run Fine-Tuning ---
trainer.train()
trainer.save_model(self.output_dir)

```

### Code Explanation:

- This code defines a standard fine-tuning loop. The `TrainingArguments` (Chunk 2.2) define the parameters (e.g., `learning_rate=2e-5`), and the `Trainer` (Chunk 2.3) handles the entire re-training process, saving the new, specialized model to `self.output_dir`.
- 

### Chunk 3: Aspect-Based Sentiment

**Flow:** Finally, the `finbert_tuner.py` module defines a function for **Aspect-Based Sentiment**. Instead of one score for an entire article, this flow allows us to ask more specific questions. For example, `get_aspect_sentiment(text, 'profit')` will scan the text *only* for profit-related keywords (e.g., 'earnings', 'revenue') and calculate a sentiment score for just that aspect, ignoring other topics. This provides a much more granular analysis.

## Section VII. Conclusion and Future Work

### Chunk 1: Key Achievements

This project successfully achieved its primary objective: to transition the stock prediction system from a price-based **regression** model to a direction-based **classification** model.

The key finding, shown in the `task7_advanced_evaluation.py` summary report, is that **sentiment features provide a significant and measurable lift in predictive performance**.

- **Baseline (Technical-Only):** The best technical-only model (Gradient Boosting) achieved an F1 Score of **65.7%**.
- **Sentiment-Enhanced:** The best sentiment-only model (also Gradient Boosting) achieved an F1 Score of **68.8%**.

This represents a **+4.8% relative improvement in F1 Score**, validating the hypothesis that news sentiment captures market psychology not reflected in historical price data alone. The final 21-model experiment provides a robust, fast (14-second runtime), and fully automated

pipeline for end-to-end sentiment-based trading analysis.

---

## Chunk 2: Future Work

Based on the project's findings, two clear paths for future work emerge:

1. **Integrate Independent Research:** The most logical next step is to replace the simple **TextBlob** model (from Section III) with the fine-tuned **FinBERT** model (from Section VI). This would combine the project's best-performing algorithm (Gradient Boosting) with its most advanced sentiment analysis tool, likely pushing performance even higher.
2. **Explore Sequential Models:** The current models (XGBoost, SVM) are "static"—they do not consider the sequence of past days. A promising direction would be to re-introduce the LSTM/GRU architectures (from Task C.4/C.5) but adapt them for this project's *classification* task, allowing the model to learn from temporal patterns in both price and sentiment.